

Gradient Descent: An Optimization Algorithm

Introduction

- **Objective:** Minimize a cost function $J(w, b)$ by systematically adjusting parameters w and b .
- Gradient descent is widely used in machine learning for training models, including deep learning models.
- Applicable not only to linear regression but also to any function minimization problem.

Overview of Gradient Descent

- **Starting Point:**
 - Begin with an initial guess for parameters w and b (often set to zero in linear regression).
- **Iterative Process:**
 - Modify w and b step by step to reduce the cost function $J(w, b)$.
 - Continue updating parameters until reaching a minimum value of J .

Generalization to Multiple Parameters

- Gradient descent extends beyond two parameters:
 - For a cost function $J(w_1, w_2, \dots, w_n, b)$, the goal is to minimize J over all parameters.
 - The method systematically updates each w_i and b to find the optimal values.

Intuition Behind Gradient Descent

- **Visualizing the Cost Function:**
 - Imagine the cost function as a surface plot where different values of w and b correspond to different heights.
 - High points represent higher cost values, and valleys represent lower cost values.
- **Descending to the Minimum:**
 - Start at an initial point on the cost surface.
 - Look for the steepest downward direction and take a small step.
 - Repeat this process iteratively to reach a local minimum.

Properties of Gradient Descent

- **Steepest Descent:**
 - At each step, move in the direction where the function decreases the most.
 - Ensures the fastest descent to a minimum.
- **Local Minima:**
 - Some functions have multiple minima.
 - The final minimum reached depends on the initial starting position.
 - If started in a different location, gradient descent might settle in a different valley (local minimum).

Implementing Gradient Descent

- **Update Rule:**
 - Gradient descent updates parameters w and b using the formulas:

$$w := w - \alpha \frac{d}{dw} J(w, b)$$

$$b := b - \alpha \frac{d}{db} J(w, b)$$

- Here, α (learning rate) determines the step size.
- The derivative terms indicate the direction and magnitude of the adjustment.

Effect of Learning Rate α

- **Too Small α :** Very slow convergence.
- **Too Large α :** Can overshoot and diverge.
- **At the Minimum:** Gradient is zero, so updates stop.

Feature Scaling for Faster Convergence

- Features with large ranges slow convergence.
- **Methods:**
 - **Min-Max Scaling:** $x' = \frac{x - x_{min}}{x_{max} - x_{min}}$
 - **Z-score Normalization:** $x' = \frac{x - \mu}{\sigma}$

Recognizing Gradient Descent Convergence

- **Learning Curve:** Plot cost vs. iterations.
- **Flat Cost Function:** Indicates convergence.

- **Threshold ϵ :** Stop if cost change is very small.

Feature Engineering and Polynomial Regression

Feature Engineering

- **What is Feature Engineering?**
 - Creating new features by transforming or combining existing ones.
 - Helps machine learning models make better predictions.
- **Example: House Price Prediction**
 - Given two features:
 - x_1 = width (frontage) of a land plot
 - x_2 = depth of a land plot
 - Basic model:

$$f(x) = w_1x_1 + w_2x_2 + b$$

- **Better Approach:** Create a new feature x_3 , where:

$$x_3 = x_1 \times x_2$$

- New model:

$$f(x) = w_1x_1 + w_2x_2 + w_3x_3 + b$$

- x_3 represents the **area** of the land, which is more predictive.

Polynomial Regression

- **Extends Linear Regression:**
 - Instead of fitting straight lines, fits curves to the data.
- **Example:** Predicting house prices using **square footage** (x).
 - A linear model may not fit the data well.
 - **Quadratic Model:**

$$f(x) = w_1x + w_2x^2 + b$$

- **Cubic Model:**

$$f(x) = w_1x + w_2x^2 + w_3x^3 + b$$

- Higher-degree polynomials allow more flexibility.

Feature Scaling in Polynomial Regression

- **Why Important?**
 - Squaring/cubing features creates vastly different ranges.
 - Example:
 - x ranges from 1 to 1,000
 - x^2 ranges from 1 to 1,000,000
 - x^3 ranges from 1 to 1,000,000,000
 - Without scaling, gradient descent struggles.
- **Alternative Polynomial Features**
 - Instead of x^2 or x^3 , use:
 - **Square Root:** $w_1x + w_2\sqrt{x} + b$
 - **Logarithm:** $w_1x + w_2\log(x) + b$
- **Choosing Features**
 - Use cross-validation to test different feature sets.
 - More features \neq better model (avoid overfitting).

Tools for Polynomial Regression

- **Scikit-Learn:**
 - Implements polynomial regression with **few lines of code**.
 - Helps avoid reimplementing algorithms manually.
 - **Why Learn Implementation?**
 - Understanding fundamentals helps in debugging models.
 - Avoid reliance on "black-box" libraries.
-

Conclusion

- **Gradient Descent:**
 - Optimizes cost function iteratively.
 - Choice of **learning rate** is critical.
 - Feature scaling speeds up convergence.
- **Feature Engineering:**
 - Creating new features improves model accuracy.
 - Example: Using **area** in house price prediction.

- **Polynomial Regression:**
 - Extends linear regression to model **curved** data.
 - Feature scaling is crucial for stability.
- **Machine Learning Practice:**
 - Implement manually to understand the math.
 - Use tools like **Scikit-Learn** for efficiency.

This version includes all the key points on **Feature Engineering and Polynomial Regression**. Let me know if you need any refinements! 🚀